


# CS188 Announcements –

---

- P0 is due Friday (ungraded, but pulse check)
  - No slip days (for this project only)
- Homework 1
  - Released now, due Sunday (no slip days!)
  - Some concepts (including “graph search” starting in Q2) covered tomorrow
- Project 1
  - Released tomorrow, due Tuesday
  - Work in teams of 1-2
- Office hours start today
- Discussion sections start tomorrow and Thursday
- Exam prep sections start Monday and Tuesday

# CS 188: Artificial Intelligence

## Search




Instructors: Anwar Baroudi and Daniel Fried

University of California, Berkeley

[These slides adapted from Dan Klein, Pieter Abbeel, Anca Dragan, and Sergey Levine]



# Today

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search




# Why should agents plan?

---



# Reflex Agents

- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS
- *Can a reflex agent be rational?*




[Demo: reflex optimal (L2D1)]

[Demo: reflex odd (L2D2)]

# Planning Agents

- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - **Consider how the world WOULD BE**
- Optimal vs. complete planning
- Planning vs. replanning



[Demo: re-planning (L2D3)]

[Demo: mastermind (L2D4)]

# Search Problems


---




# Search Problems

- A **search problem** consists of:

- A state space



- A successor function  
(with actions, costs)




- A start state and a goal test
- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

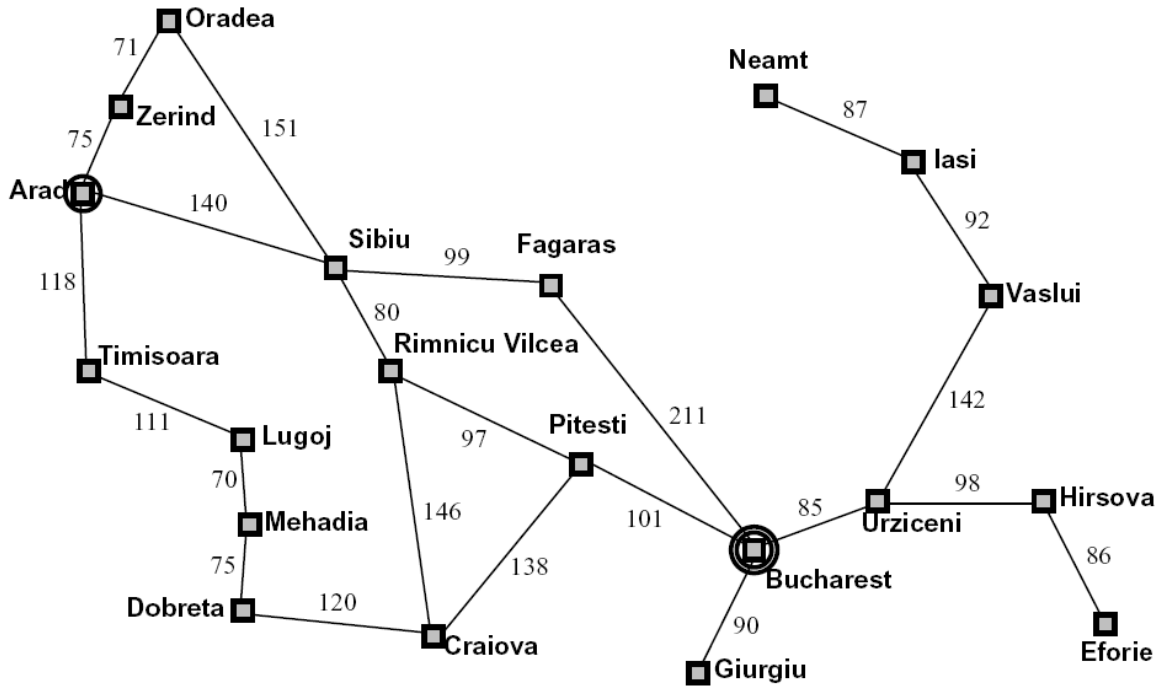


# Search Problems Are Models

---




# Example: Traveling in Romania



- State space:
  - Cities
- Successor function:
  - Roads: Go to adjacent city with cost = distance
- Start state:
  - Arad
- Goal test:
  - Is state == Bucharest?
- Solution?

# What's in a State Space?

The **world state** includes every last detail of the environment




A **search state** keeps only the details needed for planning (abstraction)


- Problem: Pathing
  - States:  $(x,y)$  location
  - Actions: NSEW
  - Successor: update location only
  - Goal test:  $\text{is } (x,y) = \text{END}$
- Problem: Eat-All-Dots
  - States:  $\{(x,y), \text{dot booleans}\}$
  - Actions: NSEW
  - Successor: update location and possibly a dot boolean
  - Goal test: dots all false

# State Space Sizes?

- World state:
  - Agent positions: 120
  - Food count: 30
  - Ghost positions: 12
  - Agent facing: NSEW
- How many
  - World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
  - States for pathing?  
120
  - States for eat-all-dots?  
 $120 \times (2^{30})$




# Safe Passage



- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)


# State Space Graphs and Search Trees

---




# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are states in the state space (abstract world configurations)
  - Edges represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



# Search Trees




- A search tree:
  - A “what if” tree of plans and their outcomes
  - The start state is the root node
  - Children correspond to successors
  - Nodes show states, but correspond to PLANS that achieve those states
  - **For most problems, we can never actually build the whole tree**



# State Space Graphs vs. Search Trees


## State Space Graph



*Each NODE in in the search tree is an entire PATH in the state space graph.*

*We construct both on demand – and we construct as little as possible.*


## Search Tree



# State Space Graphs vs. Search Trees

---

Consider this 4-state graph:





How big is its search tree (from S)?



# State Space Graphs vs. Search Trees

Consider this 4-state graph:


## How big is its search tree (from S)?




Important: Lots of repeated structure in the search tree!

# Tree Search


---



# Search Example: Romania



# Searching with a Search Tree



- Search:

- Expand out potential plans (one node per plan)
- Maintain a **fringe** of partial plans under consideration (nodes in white)
- Try to expand as few tree nodes as possible

# General Tree Search


---

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
  - Fringe (contains candidates)
  - Expansion (uses successor function)
  - Exploration strategy (defines the search algorithm)
- Main question: which fringe nodes to explore?


# Example: Tree Search

---





## Search tree



```

graph TD
    S((S)) --- d((d))
    S --- e1((e))
    S --- p1((p))
    d --- b((b))
    d --- c((c))
    d --- e2((e))
    e1 --- h1((h))
    e1 --- r1((r))
    p1 --- q1((q))
    b --- a1((a))
    c --- a2((a))
    e2 --- h2((h))
    e2 --- r2((r))
    h1 --- p2((p))
    h1 --- q2((q))
    r1 --- f1((f))
    r1 --- G1((G))
    p2 --- q3((q))
    q2 --- c1((c))
    q2 --- G2((G))
    f1 --- c2((c))
    f1 --- G3((G))
    c1 --- a3((a))
    c2 --- a4((a))
    G1 --- a5((a))
    G2 --- a6((a))
    G3 --- a7((a))
    G4((G)) --- a8((a))
  
```

~~s~~

~~s → d~~

s → e

s → p

s → d → b

s → d → c

~~s → d → e~~

s → d → e → h

~~s → d → e → r~~

~~s → d → e → r → f~~

s → d → e → r → f → c

~~s → d → e → r → f → G~~

# Depth-First Search



---



# Depth-First Search


*Strategy: expand a  
deepest node first*

*Implementation:  
Fringe is a LIFO stack*



# Search Algorithm Properties

---



# Search Algorithm Properties


- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?

- Cartoon of search tree:

- $b$  is the branching factor (max number of successors)
- $m$  is the maximum depth
- solutions at various depths


- Number of nodes in entire tree?

- $1 + b + b^2 + \dots + b^m = O(b^m)$




# Depth-First Search (DFS) Properties

- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the fringe take?
  - Only has siblings on path to root, so  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
  - No, it finds the “leftmost” solution, regardless of depth or cost



# Breadth-First Search



---



# Breadth-First Search

*Strategy: expand a shallowest node first*


*Implementation: Fringe is a FIFO queue*






# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Processes all nodes above shallowest solution
  - Let depth of shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, so yes!
- Is it optimal?
  - Only if all actions have the same cost (more on costs later)



# Quiz: DFS vs BFS

---




# DFS vs BFS

---


- When will BFS outperform DFS?
- When will DFS outperform BFS?

# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. ....
- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!



# Cost-Sensitive Search




BFS finds the shortest path in terms of number of actions.  
It does not find the least-cost path. We will now cover  
a similar algorithm which does find the least-cost path.

How?

# Uniform Cost Search



---



# Uniform Cost Search


Strategy: expand a  
cheapest node first:

Fringe is a priority queue  
(priority: cumulative cost)



# Uniform Cost Search (UCS) Properties



- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the “effective depth” is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\varepsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof next lecture via  $A^*$ )





# Uniform Cost Issues


- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location
- We'll fix that soon!



[Demo: maze with deep/shallow water DFS/BFS/UCS (L2D7)]

# The One Queue

- All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you can avoid the  $\log(n)$  overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object



# Search and Models

---

- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all “in simulation”
  - Your search is only as good as your models...

